

Specification and Implementation of Distributed Systems

Seif Haridi - Royal Institute of Technology
Peter Van Roy - Université catholique de Louvain
haridi(at)kth.se
peter.vanroy(at)uclouvain.be


9/27/21

1

■ 1

Need for Distributed Abstractions

- Core of any distributed system is a set of distributed algorithms
 - Implemented as a middleware between network (OS) and the application
- **Reliable** applications need underlying services **stronger** than network protocols (eg TCP, UDP)



9/27/21

2

■ 2

Need for Distributed Abstractions (2)

- Network protocols aren't enough

- Communication

- Reliability guarantees (eg TCP) only offered for **one-to-one** communication (client-server)
 - How to do group communication?

- High-level services

- Sometimes one-to-many communication isn't enough
 - Need reliable high-level services

Abstractions in this course

reliable broadcast
causal order broadcast
total order broadcast
terminating reliable broadcast

shared memory
consensus
atomic commit
group membership

9/27/21

3

■3

Reliable Distributed Abstractions

- Example 1: *reliable broadcast*

- Ensure that a message sent to a group of nodes is received (delivered) by **all or none**

- Example 2: *atomic commit*

- Ensure that the nodes reach the **same** decision on whether to commit or abort a transaction

9/27/21

4

■4

Event-based Component Model

9/27/21

5

■5

The Model

- Each node models a sequential program
- Assume a global clock
- At each tick, one node takes one step
 - Computation Step: local computation
 - Communication Step: deliver a message
- Different models
 - Receive 1 msg and send 1 msg [Guerraoui]
 - At most receive 1 msg, and send at most 1 msg to each neighbor [Lynch]
 - Receive k msgs, and send at most 1 msg to each neighbor [Attiya & Welch]
- We saw Attiya & Welch, but the others are very similar

9/27/21

6

■6

Event-based Programming

- Nodes execute programs:

- Each program consists of a set of **modules** or **component specifications**.
- At runtime these are deployed as components
- The components form a software stack
- Components interact via events (with attributes):

```
upon event <ReceiveEvent, attr1, attr2,...> do
    // local computation (single-threaded)
    ...
    trigger <SendEvent, attr3, attr4,...>
```

9/27/21

7

■7

Event-based Programming

- Events give different kinds of information

- Messages (most of the time)
- Timers
- Boolean conditions (e.g. $x==5 \ \& \ y<9$)

- Three event categories

- **Requests** (flow downward)
- **Indications** (the responses, flow upward)
- **Confirmations** (special type of indication, like an OK or ACK)

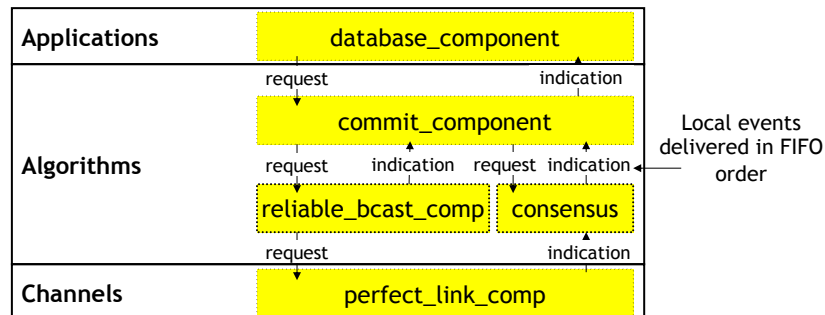
9/27/21

8

■8

Modules on one Node

- Stack of **modules** on a single node
- Requests flow down, indications flow up



9/27/21

9

■9

Send and Receive Instructions

- Send instruction is called **trigger**:
 - Send a request or an indication
 - **Example**: Send a message with data
- Receive instruction is called **upon event**:
 - Receive a request or an indication
 - **Example**: Receive a message with data

```
trigger <sendBcast | dest, [data1, data2, ...] >
```

```
upon event <delBcast | src, [data1,data2, ...]> do
```

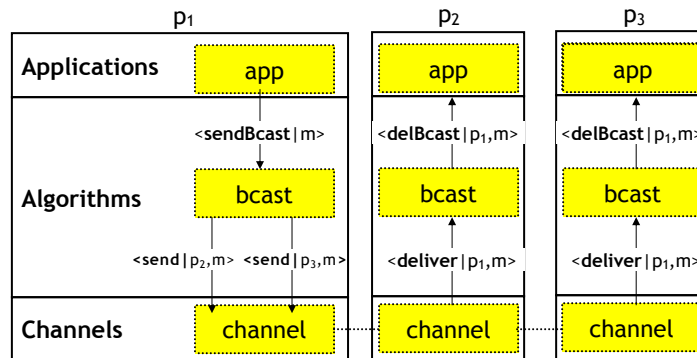
9/27/21

10

■10

Modules on several Nodes

- This application uses a Broadcast component
 - Which uses a channel module to broadcast



9/27/21

11

■ 11

Specification

9/27/21

12

■ 12

Specification of a Service

■ How to specify a distributed service?

□ Interface (a.k.a. Contract, API)

- Requests
- Responses

□ Correctness Properties

- Safety
- Liveness

□ Model

- Failure assumptions
- Synchrony assumptions

□ Implementation

- Composed of other services
- Adheres to interface and satisfies correctness
- Has internal events

declarative
specification
“what”
a.k.a. **problem**

imperative,
many possible
“how”

9/27/21

13

■13

Simple Interface Example (1)

Module 1.1 Interface of a printing module

Module:

Name: Print.

Events:

Request: $\langle \text{PrintRequest} \mid \text{rqid}, \text{str} \rangle$: Requests a string to be printed. The token rqid is an identifier of the request.

Confirmation: $\langle \text{PrintConfirm} \mid \text{rqid} \rangle$: Used to confirm that the printing request with identifier rqid succeeded.

9/27/21

14

■14

Simple Interface Example (2)

Algorithm 1.1 Printing service

Implements:

Print.

```
upon event  $\langle \textit{PrintRequest} \mid \textit{rqid}, \textit{str} \rangle$  do
  print str;
  trigger  $\langle \textit{PrintConfirm} \mid \textit{rqid} \rangle$ ;
```

9/27/21

Seif Haridi, ID2203

15

■15

Simple Interface Example (3)

Module 1.2 Interface of a bounded printing module

Module:

Name: BoundedPrint.

Events:

Request: $\langle \textit{BoundedPrintRequest} \mid \textit{rqid}, \textit{str} \rangle$: Request a string to be printed. The token *rqid* is an identifier of the request.

Confirmation: $\langle \textit{PrintStatus} \mid \textit{rqid}, \textit{status} \rangle$: Used to return the outcome of the printing request: Ok or Nok.

Indication: $\langle \textit{PrintAlarm} \rangle$: Used to indicate that the threshold was reached.

9/27/21

Seif Haridi, ID2203

16

■16

Simple Interface Example (4)

Algorithm 1.2 Bounded printer based on (unbounded) printing service

Implements:

BoundedPrint.

Uses:

Print.

```
upon event < Init > do
  bound := PredefinedThreshold;

upon event < BoundedPrintRequest | rqid, str > do
  if bound > 0 then
    bound := bound-1;
    trigger < PrintRequest | rqid, str >;
    if bound = 0 then trigger < PrintAlarm >;
  else
    trigger < PrintStatus | rqid, Nok >;

upon event < PrintConfirm | rqid > do
  trigger < PrintStatus | rqid, Ok >;
```

<Init> automatically
generated upon
component creation

9/27/21

Seif Haridi, ID2203

17

■17

Safety and Liveness

- Correctness is always expressed in terms of
 - Safety and liveness
- Safety
 - Properties that state that “something bad **never** happens”
- Liveness
 - Properties that state that “something good **eventually** happens”

9/27/21

Seif Haridi, ID2203

18

■18

Correctness Example (1)

- Correctness of **You** in LSINF2345

- Safety

- You should **never** fail the exam

- Liveness

- You should **eventually** take the exam

9/27/21

Seif Haridi, ID2203

19


■19

Correctness Example (2)

- Correctness of traffic lights at intersection

- Safety

- More than one direction should **never** have a green light at same time



- Liveness

- Every direction should **eventually** get a green light

9/27/21

Seif Haridi, ID2203

20

■20

Safety & Liveness All That Matters

- A **property** P is a function that takes an execution (!) and returns true/false
 - I.e., P is a **predicate**
- “Any [property] can be expressed as the conjunction of a **safety** property and a **liveness** property”
 - Alpern & Schneider, Inf.Proc.Letters 1985

9/27/21

Seif Haridi, ID2203

21

■21

Prefixes and Extensions

- The **prefix** of execution E is the first k (for some $k > 0$) configurations of E , with events
 - I.e., cut off the tail of E
 - I.e., finite beginning of E
- An **extension** of a prefix P is any execution that has P as a prefix
 - The extension continues P

9/27/21

Seif Haridi, ID2203

22

■22

Safety Formally Defined

- Informally, property P is a safety property if
 - Every execution E violating P “goes bad”, i.e., it has a bad event s.t. every execution starting like E and behaving like E up to the bad event (including), will violate P regardless of what it does afterwards
 - When an execution “goes bad”, it “stays bad”!
- Formally, a property P is a **safety** property if
 - Given any execution E such that $P(E)=\text{false}$,
 - There exists a prefix of E s.t. every extension of that prefix gives an execution F s.t. $P(F)=\text{false}$

9/27/21

Seif Haridi

23

■23

Safety Example

- Point-to-point message communication
 - Safety P :
 - A message sent is delivered **at most** once
 - Take an execution where a message is delivered more than once
 - Cut off the tail after the second delivery
 - Any extension will give an execution which also violates the required property

9/27/21

Seif Haridi, ID2203

24

■24

Liveness Formally Defined

- A property P is a **liveness** property if
 - Given any execution E such that $P(E)=\text{true}$,
 - For every prefix F of E , there exists an extension of F s.t. $P(F)=\text{true}$
- “As long as there is life there is hope”

9/27/21

Seif Haridi, ID2203

25

■25

Liveness Example

- Point-to-point message communication
 - Liveness P :
 - A message sent is delivered **at least** once
 - Take the prefix of any execution
 - If prefix contains delivery, any extension satisfies P
 - If prefix doesn't contain the delivery, extend it so that it contains a delivery, then the prefix + extended part will satisfy P


9/27/21

Seif Haridi, ID2203


26

■26

Formal Definitions Visually



- Safety can always be made false in finite time
- **Safety is false** for an execution E then there exists a prefix such that all extensions are false (system is broken)



- Liveness can only be made false in infinite time
- **Liveness is true** for an execution E then for all prefixes there exists an extension that is true (system can progress)

9/27/21

27

■ 27

More on Safety

- Safety can only be
 - **satisfied** in infinite time (you're never safe)
 - **violated** in finite time (when the bad happens)
- Often involves the word "never", "at most", "cannot",...
- Sometimes called "partial correctness"

9/27/21

Seif Haridi, ID2203

28

■ 28

More on Liveness

- Liveness can only be
 - **satisfied** in finite time (when the good happens)
 - **violated** in infinite time (there's always hope)
- Often involves the words “eventually”, “must”, “at least”
 - Eventually means at some (often unknown) point in “future”
- Liveness is often just “termination”

9/27/21

Seif Haridi, ID2203

29

■29

Pondering Safety and Liveness

- Why not define safety to be a predicate that is true in every configuration in an execution? **[d]**
- Is every property really either liveness or safety?
 - Every message should be delivered exactly 1 time **[d]**

9/27/21

Seif Haridi, ID2203

30

■30

Failure Assumptions for Nodes

9/27/21

31

■ 31

Model/Assumptions

- Specification needs to specify the model
 - Assumptions needed for the algorithm to be correct
- Model includes **assumptions** on
 - Failure behavior of processes & channels
 - Timing behavior of processes & channels

9/27/21

Seif Haridi, ID2203

32

■ 32

Node Failures

- Nodes may fail in four ways:
 - Crash-stop
 - Omission
 - Crash-recovery
 - Byzantine/Arbitrary
- Nodes that don't fail in an execution are correct

9/27/21

Seif Haridi, ID2203

33

■33

Crash-stop Failures

- Crash-stop failure
 - Node stops taking steps
 - Not sending msgs
 - Nor receiving msgs
- Default failure model is crash-stop
 - Hence, do not recover
 - But are nodes really not allowed to recover? [d]

9/27/21

Seif Haridi, ID2203

34

■34

Omission Failures

- Node omits sending or receiving messages
 - Some differentiate between
 - Send omission
 - Not sending msgs node has to send according to its algo
 - Formally, an event removing element from outbuf[i]
 - Receive omission
 - Not receiving messages that have been sent to node
 - Formally, an event removing element from inbuf[i]
 - For us, omission failure covers both types...

9/27/21

Seif Haridi, ID2203

35

■35

Crash-recovery Failures

- The node might crash
 - It stops taking steps, not receiving and sending messages
- It may **recover** after crashing
 - Special **<Recovery>** event automatically generated
 - Formal model: Restarting in some **initial recovery state**
- Has access to **stable storage**
 - May read/write (**expensive**) to permanent storage device
 - Storage survives crashes
 - E.g., save state to storage, crash, recover, read saved state...

9/27/21

Seif Haridi, ID2203

36

■36

Crash-recovery Failures (2)

- Failure different in crash-recovery model
 - A node is **faulty** in an execution if
 - It crashes and never recovers, or
 - It crashes and recovers infinitely often (**unstable**)
 - Hence, a **correct node** may crash and recover
 - As long as it is a finite number of times
 - A correct node will eventually (after finite time) do no more crashes

9/27/21

Seif Haridi, ID2203

37

■37

Byzantine Failures

- Byzantine/Arbitrary failures
 - A node may behave arbitrarily
 - Sending messages not specified by its algorithm
 - Updating its state not specified by its algorithm
 - May behave **maliciously**, attacking the system
 - Several malicious nodes might collude (work together to break the system)

9/27/21

Seif Haridi, ID2203

38

■38

Fault-tolerance Hierarchy

- There is a hierarchy among the failure types
 - Which one is a special case of which? [d]
- Crash-stop is a special case of Omission
 - Omission restricted to omitting everything after a certain event

9/27/21

Seif Haridi, ID2203

39

■39

Fault-tolerance Hierarchy (2)

- Assume a special case of crash-recovery
 - Where nodes use only stable storage (no other memory)
- In this case crash-recovery is identical to omission
 - Crashing, recovering, and reading last state from storage
 - Just same as omitting send/recv while being crashed
- Assume a special case of crash-recovery
 - Where nodes use some volatile memory
 - Then recovered nodes might not be able to restore all of state
 - Thus crash-recovery extends omission with **amnesia**
- Therefore Omission is a special case of Crash-recovery
 - Omission restricted to omitting everything after a certain event
 - Possibly recovering, not allowing for amnesia

9/27/21


Seif Haridi, ID2203

40

■40

Fault-tolerance Hierarchy (3)

- Crash-recovery is a special case of Byzantine
 - Since Byzantine allows anything
- Byzantine tolerance \rightarrow crash-recovery tol.
 - Crash-recovery \rightarrow omission, omission \rightarrow crash



9/27/21

Seif Haridi, ID2203

41

■41

Failure Assumptions for Channels

9/27/21

42

■42

Channel failure modes

■ Fair-Loss Links

- Channel delivers any message sent with non-zero probability (no network partitions)

■ Stubborn Links

- Channel delivers any message sent infinitely many times

■ Perfect Links

- Channel that delivers any message sent exactly once


9/27/21

Seif Haridi, ID2203

43

■43

Fair Loss Links



9/27/21

Seif Haridi, ID2203

44

■44

Fair-loss links: Interfaces

- **Module:**

- Name: FairLossPointToPoint (flp2p)

- **Events:**

- **Request:** $\langle \text{flp2pSend} \mid \text{dest}, m \rangle$
 - Request transmission of message m to node dest
- **Indication:** $\langle \text{flp2pDeliver} \mid \text{src}, m \rangle$
 - Deliver message m sent by node src

- **Properties:**

- *FL1, FL2, FL3.*

9/27/21

Seif Haridi, ID2203

45

■45

Fair-loss links

- **Properties**

- ***FL1. Fair-loss:*** If m is sent infinitely often by p_i to p_j , and neither crash, then m is delivered infinitely often by p_j
- ***FL2. Finite duplication:*** If a m is sent a finite number of times by p_i to p_j , then it is delivered a finite number of times by p_j
 - I.e., a message cannot be duplicated infinitely many times
- ***FL3. No creation:*** No message is delivered unless it was sent

9/27/21

Seif Haridi, ID2203

46

■46

Stubborn links: interface

■ Module:

- Name: StubbornPointToPoint (sp2p)

■ Events:

- Request: $\langle \text{sp2pSend} \mid \text{dest}, m \rangle$
 - Request the transmission of message m to node dest
- Indication: $\langle \text{sp2pDeliver src}, m \rangle$
 - deliver message m sent by node src

■ Properties:

- *SL 1, SL 2*

9/27/21

Seif Haridi, ID2203

47

■47

Stubborn links

■ Properties

- *SL 1. Stubborn delivery:* if a node p_i sends a message m to a correct node p_j , and p_i does not crash, then p_j delivers m an infinite number of times
- *SL 2. No creation:* if a message m is delivered by some node p_j , then m was previously sent by some node p_i

9/27/21

Seif Haridi, ID2203

48

■48

Implementing Stubborn Links

■ Implementation

- Use the Lossy link (fair-loss link)
- Sender stores every message it sends in **sent**
- It periodically resends all messages in **sent**

■ Correctness

- **SL1. Stubborn delivery**
 - If node doesn't crash, it will send every message infinitely many times. Messages will be delivered infinitely many times. Lossy link may only drop a (large) fraction.
- **SL2. No creation**
 - Guaranteed by the Lossy link

9/27/21

Seif Haridi, ID2203

49

■49

Stubborn Links in the Formal Model

■ Consider an execution

$$E = \langle c_0, e_1, c_1, e_2, c_2, \dots \rangle$$

■ Stubborn property

- If message m appears in $\text{outbuf}[p]$ for link $i \leftrightarrow j$ of a node i in E
 - Then $\text{del}'(i, j, m)$ will happen infinitely often, without removing m from $\text{outbuf}[p]$

■ Finite duplication and No creation

- Similarly

9/27/21

Seif Haridi, ID2203

50

■50

Algorithm (sl)

- **Implements:** StubbornLinks (sp2p)
- **Uses:** FairLossLinks (flp2p)
- **upon event** $\langle \text{Init} \rangle$ **do**
 - $\text{sent} := \emptyset$
 - $\text{startTimer}(\text{TimeDelay})$
- **upon event** $\langle \text{Timeout} \rangle$ **do**
 - **forall** $(\text{dest}, m) \in \text{sent}$ **do**
 - **trigger** $\langle \text{flp2pSend} \mid \text{dest}, m \rangle$
 - $\text{startTimer}(\text{TimeDelay})$

9/27/21

Seif Haridi, ID2203

51

■51

Algorithm (sl) (2)

- **upon event** $\langle \text{sp2pSend} \mid \text{dest}, m \rangle$ **do**
 - $\text{sent} := \text{sent} \cup \{ (\text{dest}, m) \}$
- **upon event** $\langle \text{flp2pDeliver} \mid \text{src}, m \rangle$ **do**
 - **trigger** $\langle \text{sp2pDeliver src}, m \rangle$

9/27/21

Seif Haridi, ID2203

52

■52

Perfect links: interface

■ Module:

- Name: PerfectPointToPoint (pp2p)

■ Events:

- Request: $\langle \text{pp2pSend} \mid \text{dest}, m \rangle$
 - Request the transmission of message m to node dest
- Indication: $\langle \text{pp2pDeliver} \mid \text{src}, m \rangle$
 - deliver message m sent by node src

■ Properties:

- *PL1, PL2, PL3*

9/27/21

Seif Haridi, ID2203

53

■53

Perfect links aka Reliable links

■ *Properties*

- *PL1. Reliable Delivery:* If neither p_i nor p_j crashes, then every message sent by p_i to p_j is eventually delivered by p_j
- *PL2. No duplication:* Every message is delivered at most once
- *PL3. No creation:* No message is delivered unless it was sent

9/27/21

Seif Haridi, ID2203

54

■54

Perfect links aka Reliable links

- Which one is safety/liveness/neither?
- **PL1. Reliable Delivery:** If neither p_i nor p_j crashes, then every message sent by p_i to p_j is eventually delivered by p_j (**liveness**)
- **PL2. No duplication:** Every message is delivered at most once (**safety**)
- **PL3. No creation:** No message is delivered unless it was sent (**safety**)

9/27/21

Seif Haridi, ID2203

55

■55

Perfect Link Implementation

- **Implementation**
 - Use Stubborn links
 - Receiver keeps **log** of all received messages in **Delivered**
 - Only deliver (pp2pDeliver) messages that weren't delivered before
- **Correctness**
 - **PL1. Reliable Delivery**
 - Guaranteed by Stubborn link. In fact the Stubborn link will deliver it infinite number of times
 - **PL2. No duplication**
 - Guaranteed by our log mechanism
 - **PL3. No creation**
 - Guaranteed by Stubborn link (and its lossy link? **[D]**)

9/27/21

Seif Haridi, ID2203

56

■56

Algorithm (pl)

- **Implements:** PerfectLinks (pp2p).
- **Uses:** StubbornLinks (sp2p).
- **upon event** $\langle \text{Init} \rangle$ **do** $\text{delivered} := \emptyset$
- **upon event** $\langle \text{pp2pSend} \mid \text{dest}, m \rangle$ **do**
 - **trigger** $\langle \text{sp2pSend dest}, m \rangle$
- **upon event** $\langle \text{sp2pDeliver} \mid \text{src}, m \rangle$ **do**
 - **if** $m \notin \text{delivered}$ **then**
 - $\text{delivered} := \text{delivered} \cup \{ m \}$
 - **trigger** $\langle \text{pp2pDeliver} \mid \text{src}, m \rangle$

9/27/21

Seif Haridi, ID2203

57

■57

Default Assumptions in Course

- We assume perfect links (aka reliable) most of the course (unless specified otherwise)
- Roughly, reliable links ensure messages exchanged between correct nodes are delivered exactly once
- NB. Messages are **uniquely** identified and
 - the message identifier includes the sender's identifier
 - i.e. if "same" message sent twice, it's considered as two different messages
- Stubborn links used mostly in the crash-recovery

9/27/21

Seif Haridi, ID2203

58

■58

TCP vs perfect links

- How does TCP efficiently maintain **delivered** log? [d]
- Use sequence numbers
 - ACK: “I have received everything up to byte X”

9/27/21

59

■59

Synchrony Assumptions

9/27/21

60

■60

Synchrony Assumptions

- Synchrony assumptions relate to
 - different processing speeds of nodes
 - different speeds of messages (channels)
- Three basic types of systems:
 - Asynchronous system
 - Synchronous system
 - Partially synchronous system

9/27/21

Seif Haridi, ID2203

61

■61

Asynchronous Systems

- No timing assumption on nodes and channels
 - Processing time varies arbitrarily
 - No bound on transmission time
- We know it well from our formal model!
 - Causality
 - Lamport clocks (or vector clocks) to observe causality
 - Total order not observable, no access to clock
 - We used Computation Theorem
- Is Internet asynchronous? [d]

9/27/21

Seif Haridi, ID2203

62

■62

Synchronous Systems

- Model assumes
 - Synchronous computation
 - Known upper bound on node processing delays
 - Synchronous communication
 - Known upper bound on message transmission delay
 - Synchronous physical clocks
 - Nodes have local physical clock
 - Known upper bound clock drift and clock skew
- Why study synchronous systems? [d]

9/27/21

Seif Haridi, ID2203

63

■63

Partial Synchrony

- Asynchronous system
 - Which **eventually** becomes synchronous
 - Cannot know when, but in every execution, some bounds eventually will hold
- It's just a way to formalize the following
 - *Your algorithm will have a long enough time window, where everything behaves nicely (synchrony), so that it can achieve its goal*
- Are there such systems? [d]

9/27/21


Seif Haridi, ID2203

64

■64

Partial Synchrony (2)

- Your algorithm will have a long *enough time* window, where everything behaves nicely (synchrony), so that it can achieve its *goal*
 - Mostly useful for *terminating* algorithms



9/27/21

Seif Haridi, ID2203

65