

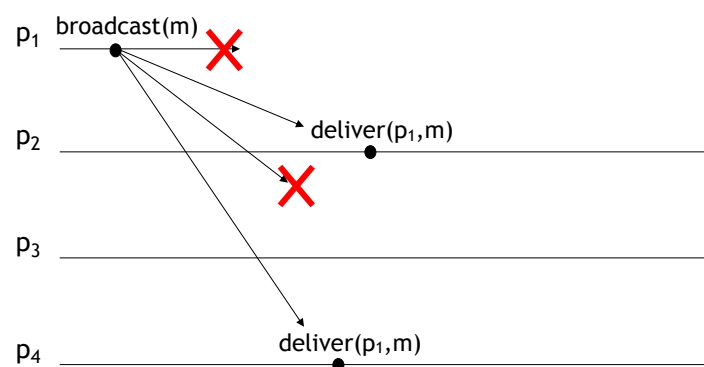
# Reliable Broadcast

Seif Haridi - Royal Institute of Technology  
Peter Van Roy - Université catholique de Louvain

haridi(at)kth.se  
peter.vanroy(at)uclouvain.be

■1

# Unreliable Broadcast



10/11/21

Seif Haridi

2

■2

## Reliable Broadcast Abstractions

- **Best-effort broadcast**
  - Guarantees reliability **only if sender is correct**
- **Reliable broadcast**
  - Guarantees reliability **independent of whether sender is correct**
- **Uniform reliable broadcast**
  - Also **considers behavior of failed nodes**
- **FIFO reliable broadcast**
  - Reliable broadcast with **FIFO delivery order**
- **Causal reliable broadcast**
  - Reliable broadcast with **causal delivery order**

10/11/21

Seif Haridi

3

■3

## Specification of Broadcast Abstractions

■4

## Best-effort broadcast (beb)

### ■ *Events*

- Request:  $\langle \text{bebBroadcast} \mid m \rangle$
- Indication:  $\langle \text{bebDeliver} \mid \text{src}, m \rangle$

### ■ *Properties: BEB1, BEB2, BEB3*

10/11/21

Seif Haridi

5

■5

## Best-effort broadcast (beb)

### ■ *Intuitively:* everything perfect unless sender crash

### ■ *Properties*

- **BEB1. Best-effort-Validity:** If  $p_i$  and  $p_j$  are **correct**, then any broadcast by  $p_i$  is eventually delivered by  $p_j$
- **BEB2. No duplication:** No message delivered more than once
- **BEB3. No creation:** No message delivered unless broadcast

10/11/21

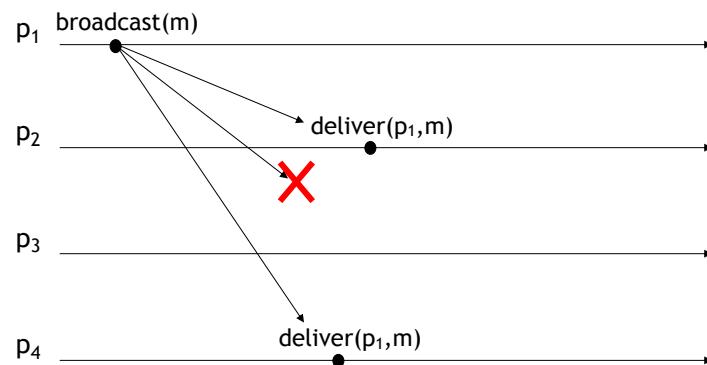
Seif Haridi

6

■6

## BEB Example

■ Is this allowed? **No**



10/11/21

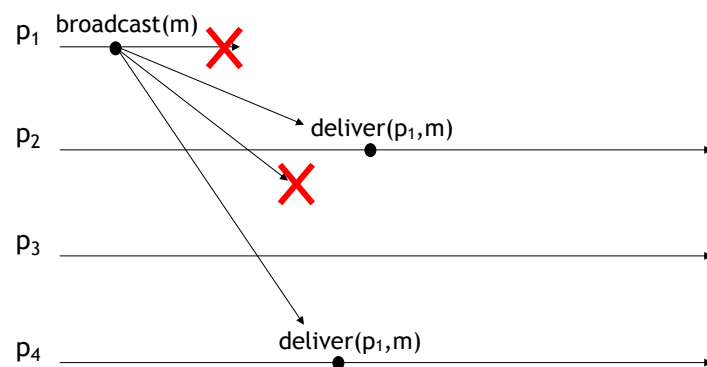
Seif Haridi

7

■7

## BEB Example (2)

■ Is this allowed? **Yes**



10/11/21

Seif Haridi

8

■8

## Reliable Broadcast

- BEB gives no guarantees if **sender crashes**
  - Strengthen to give guarantees if sender crashes
- Reliable Broadcast Intuition
  - Same as BEB, plus
  - If sender crashes:  
ensure *all or none* of the correct nodes get msg

10/11/21

Seif Haridi

9

■9

## Reliable Broadcast (rb)

- **Events**
  - Request:  $\langle \text{rbBroadcast} \mid m \rangle$
  - Indication:  $\langle \text{rbDeliver} \mid \text{src}, m \rangle$
- **Properties: *RB1, RB2, RB3, RB4***

10/11/21

Seif Haridi

10

■10

# Reliable Broadcast Properties

## ■ *Properties*

- *RB1 = BEB1. Validity*
- *RB2 = BEB2. No duplication*
- *RB3 = BEB3. No creation*
- *RB4. Agreement.*
  - If a **correct node delivers** *m*, then every correct node delivers *m*

10/11/21

Seif Haridi

11

■11

# Refining correctness

## ■ Can weaken RB1 without any effect

### Old Validity

←equivalent with→

### New Validity

- *RB1 = BEB1 Validity*
  - If *p<sub>i</sub>* and *p<sub>j</sub>* are **correct**, then any broadcast by *p<sub>i</sub>* is eventually delivered by *p<sub>j</sub>*
- *RB2 = BEB2. No duplication*
- *RB3 = BEB3. No creation*
- *RB4. Agreement.*
  - If a **correct node delivers** *m*, then every correct node delivers *m*

- *RB1 Validity.*
  - If **correct** *p<sub>i</sub>* broadcasts *m*, *p<sub>i</sub>* itself eventually delivers *m*
- *RB2 = BEB2. No duplication*
- *RB3 = BEB3. No creation*
- *RB4. Agreement.*
  - If a **correct node delivers** *m*, then every correct node delivers *m*

10/11/21


Seif Haridi

12

■12

## RB Example

- Is this allowed? **Yes**



10/11/21


Seif Haridi

13

■13

## RB Example

- Is this allowed? **Yes**



10/11/21


Seif Haridi

14

■14

## RB Example

- Is this allowed? **No**



10/11/21

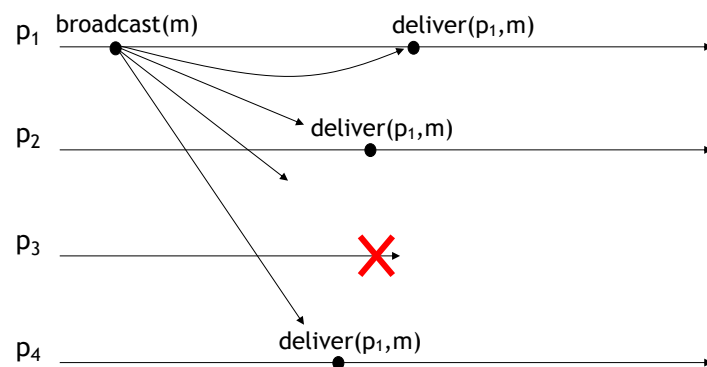
Seif Haridi

15

■15

## RB Example

- Is this allowed? **Yes**



10/11/21

Seif Haridi

16

■16



## Uniform Reliable Broadcast

- Assume the broadcast enforces some real-world action
  - Printing a message on paper
  - Withdrawing money from account in variable
  - Launching a missile
- Assume sender broadcasts message
  - Sender fails
  - No correct node delivers message
  - *Failed nodes* might or might not deliver message, is it ok?
- **Uniform** reliable broadcast intuition
  - If a *failed node* delivers, *everyone* must deliver...  
(At least the correct nodes, we cannot revive the dead...)

10/11/21

Seif Haridi

17

■17

## Uniform Reliable Broadcast (urb)

- **Events**
  - Request:  $\langle \text{urbBroadcast} \mid m \rangle$
  - Indication:  $\langle \text{urbDeliver} \mid \text{src}, m \rangle$
- **Properties:**
  - **URB1**
  - **URB2**
  - **URB3**
  - **URB4**

10/11/21

Seif Haridi

18

■18

## Uniform Broadcast Properties

### ■ *Properties*

- $URB1 = RB1.$
- $URB2 = RB2.$
- $URB3 = RB3.$
- **URB4. *Uniform Agreement*:** For any message  $m$ , if a process delivers  $m$ , then every correct process delivers  $m$

Wanted:  
Dead AND  
Alive!

10/11/21

Seif Haridi

19

■19

## Implementation of Broadcast Abstractions

■20

## Implementing BEB

- Use Perfect channel abstraction
  - Upon `<bebBroadcast | m>` send message to all nodes (for-loop)
- Correctness
  - If sender doesn't crash, every other correct node receives message by perfect channels
  - No creation & no duplication already guaranteed by perfect channels

10/11/21

Seif Haridi

21

■21

## Fail-Stop: Lazy Reliable Broadcast

- Requires perfect failure detector (P)
- To broadcast `m`:
  - **best-effort broadcast** `m`
  - When get `bebDeliver`
    - Save message, and
    - `rbDeliver` message
- If **sender `s`** crashes, detect & relay msgs from `s` to all
  - case 1: get `m` from `s`, detect crash `s`, redistribute `m`
  - case 2: detect crash `s`, get `m` from `s`, redistribute `m`
    - Why case 2? **[d]**
- Filter duplicate messages

10/11/21

Seif Haridi

22

■22

## Lazy Reliable Broadcast

- **Implements:** ReliableBroadcast (rb)

- **Uses:**

- BestEffortBroadcast (beb)
- PerfectFailureDetector (P)

- **upon event**  $\langle \text{Init} \rangle$  **do**

- $\text{delivered} := \emptyset$
- $\text{correct} := \Pi$
- **forall**  $p_i \in \Pi$  **do**  $\text{from}[p_i] := \emptyset$

for filtering  
duplicates

storage for saved  
messages

- **upon event**  $\langle \text{rbBroadcast} \mid m \rangle$  **do**

- **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, \text{self}, m) \rangle$

10/11/21

Seif Haridi

23

■23

## Lazy Reliable Broadcast (2)

- **upon event**  $\langle \text{crash} \mid p_i \rangle$  **do**

- $\text{correct} := \text{correct} \setminus \{p_i\}$
- **forall**  $(s_m, m) \in \text{from}[p_i]$  **do**  
    **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, s_m, m) \rangle$

Case 1: redistribute  
anything we have  
from failed node

- **upon event**  $\langle \text{bebDeliver} \mid p_i, (\text{DATA}, s_m, m) \rangle$  **do**

- **if**  $m \notin \text{delivered}$  **then**
- $\text{delivered} := \text{delivered} \cup \{m\}$
- $\text{from}[p_i] := \text{from}[p_i] \cup \{(s_m, m)\}$
- **trigger**  $\langle \text{rbDeliver} \mid s_m, m \rangle$
- **if**  $p_i \notin \text{correct}$  **then**  
    **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, s_m, m) \rangle$

Avoid duplicates

Store for future

Case 2: redistribute

10/11/21


Seif Haridi

24

■24

## RB Example

■ Which case? **Case 2**



10/11/21


Seif Haridi

25

■25

## RB Example

■ Which case? **Case 1**



10/11/21

Seif Haridi

26

■26

## Correctness of Lazy RB

- **RB1-RB3** satisfied by BEB
- Need to prove **RB4**
  - If a **correct node delivers**  $m$ , then every correct node delivers  $m$
- If correct  $p_j$  delivers msg broadcast by  $p_i$ 
  - If  $p_i$  is correct, BEB ensures correct delivery
  - If  $p_i$  crashes,
    - $p_j$  detects this (completeness)
    - $p_j$  uses BEB to ensure (BEB1) every correct node gets it
  - This is a proof by induction. Why? **[d]**

10/11/21

Seif Haridi

27

■27

## Eager Reliable Broadcast

- What happens if we replace  $P$  with  $\Diamond P$ ? **[d]**
  - Only affects performance, not correctness.
- Can we modify Lazy RB to not use  $P$ ? **[d]**
  - Just assume all nodes failed
  - BEB Broadcast as soon as you get a msg

10/11/21

Seif Haridi

28

■28

## Eager Reliable Broadcast

- **Uses:** BestEffortBroadcast (beb)
- **upon event**  $\langle \text{Init} \rangle$  **do**
  - $\text{delivered} := \emptyset$
- **upon event**  $\langle \text{rbBroadcast} \mid m \rangle$  **do**
  - $\text{delivered} := \text{delivered} \cup \{m\}$
  - **trigger**  $\langle \text{rbDeliver} \mid \text{self}, m \rangle$  ← **Immediately deliver**
  - **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, \text{self}, m) \rangle$  ← **Immediately BEB broadcast**
- **upon event**  $\langle \text{bebDeliver} \mid p_i, (\text{DATA}, s_m, m) \rangle$  **do**
  - **if**  $m \notin \text{delivered}$  **then**
  - $\text{delivered} := \text{delivered} \cup \{m\}$
  - **trigger**  $\langle \text{rbDeliver} \mid s_m, m \rangle$  ← **Immediately deliver**
  - **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, s_m, m) \rangle$  ← **Immediately BEB broadcast**

10/11/21

Seif Haridi

29

■29

## Correctness of Eager RB

- ***RB1-RB3*** satisfied by BEB
- Need to prove ***RB4***
  - If a **correct node delivers**  $m$ , then every correct node delivers  $m$
- If correct  $p_j$  delivers message bcast by  $p_i$ 
  - $p_j$  uses BEB to ensure (BEB1) every correct node gets it

10/11/21

Seif Haridi

30

■30

## Uniformity

- Is the proposed algorithm also **uniform**? [d]
- No.
  - Sender p immediately RB delivers and crashes
  - Only p delivers message
- Uniformity necessitates
  - If a **failed** node delivers a message m then **every correct** node delivers m

10/11/21

Seif Haridi

31

■31

## Uniform Eager RB

- Necessary condition for URB delivery
  - All correct nodes will get the msg
  - How do we know the correct nodes got msg? [d]
- Messages are **pending** until all correct nodes get it
  - Collect acks from nodes that got msg ← 

Use vector **ack[m]** at p:  
Set of nodes who  
acked m
- Deliver when all correct nodes acked
  - Use perfect FD
  - **function** canDeliver(m):
    - return correct  $\subseteq$  ack[m]

10/11/21

Seif Haridi

32

■32



## Uniform Eager RB implementation

- **upon event**  $\langle \text{urbBroadcast} \mid m \rangle$  **do**
  - $\text{pending} := \text{pending} \cup \{\text{self}, m\}$  ← remember sent messages
  - **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, \text{self}, m) \rangle$
  
- **upon event**  $\langle \text{bebDeliver} \mid p_i, (\text{DATA}, s_m, m) \rangle$  **do**
  - $\text{ack}[m] := \text{ack}[m] \cup \{p_i\}$  ←  $p_i$  obviously got  $m$
  - **if**  $(s_m, m) \notin \text{pending}$  **then** ← avoid resending
    - $\text{pending} := \text{pending} \cup \{(s_m, m)\}$
    - **trigger**  $\langle \text{bebBroadcast} \mid (\text{DATA}, s_m, m) \rangle$
  
- **Upon exists**  $(s_m, m) \in \text{pending}$  **s.t.**  $\text{canDeliver}(m)$  **and**  $m \notin \text{delivered}$  **do**
  - $\text{delivered} := \text{delivered} \cup \{m\}$
  - **trigger**  $\langle \text{urbDeliver} \mid s_m, m \rangle$  ← deliver when all correct nodes have acked

10/11/21

Seif Haridi

33

■33

## Correctness of Uniform RB

- **No creation** from BEB
- **No duplication** by using *delivered* set
- **Lemma**
  - If a correct node  $p_i$  bebDelivers  $m$ , then  $p_i$  eventually urbDelivers  $m$
- **Proof**
  - Correct node  $p_i$  bebBroadcasts  $m$  as soon as it gets  $m$ 
    - By BEB1 every correct node gets  $m$  and bebBroadcasts  $m$
    - $p_i$  gets bebDeliver( $m$ ) from every correct node by BEB1
    - By completeness of  $P$ , it will not wait for dead nodes forever
      - $\text{canDeliver}(m)$  becomes true and  $p_i$  delivers  $m$

10/11/21

Seif Haridi

34

■34

## Correctness of Uniform RB

### Validity

- If sender  $s$  is correct, it'll by **validity** (BEB1)  $\text{bebDeliver}(m)$
- By the **lemma**, it will eventually  $\text{urbDeliver}(m)$

### Uniform agreement

- Assume some node (possibly failed) URB delivers  $m$ 
  - Then  $\text{canDeliver}(m)$  was true,
    - by **accuracy** of  $P$  every correct node has BEB delivered  $m$
- By **lemma** each of the nodes that BEB delivered  $m$  will URB deliver  $m$


10/11/21

Seif Haridi

35

■35

## URB Eager Algorithm Example



10/11/21

Seif Haridi

36

■36

## How difficult is URB?

- Strong failure detectors necessary for URB?
  - No, we'll provide URB for **fail-silent** model
- Assume a **majority** of correct nodes
  - Majority =  $\lfloor n/2 \rfloor + 1$ , i.e. 6 of 11, 7 of 12...
- Every node eagerly BEB broadcast  $m$ 
  - URB deliver  $m$  when received  $m$  from a majority

10/11/21

Seif Haridi

37

■37

## Majority-ACK Uniform RB

- Same algorithm as uniform eager RB
  - Replace one function
    - **function canDeliver( $m$ )**
      - return  $|ack[m]| > n/2$  ← 

majority has acknowledged  $m$
- Agreement (main idea)
  - If a node URB delivers, it got ack from majority
  - In that majority, one node,  $p$ , must be correct
  - $p$  will ensure all correct nodes BEB deliver  $m$ 
    - The correct nodes (majority) will ack and URB deliver

10/11/21

Seif Haridi

38

■38

## Majority-ACK Uniform RB

### ■ Validity

- If correct sender sends  $m$ 
  - All correct nodes BEB deliver  $m$
  - All correct nodes BEB broadcast
  - Sender receives a majority of acks
  - Sender URB delivers  $m$